

```

# documents
doc1, doc2='I like dogs.', 'I hate dogs.'

# Split the documents and create tokens
doc1_tokens=set(doc1.lower().split())
doc2_tokens=set(doc2.lower().split())

#Print the tokens
print(doc1_tokens,doc2_tokens)

# Calculate the Jaccard Similarity
print(len(doc1_tokens.intersection(doc2_tokens)))
print(len(doc1_tokens.union(doc2_tokens)))
print(doc1_tokens.union(doc2_tokens))
jaccard_similarity=
len(doc1_tokens.intersection(doc2_tokens))/len(doc1_tokens.union(doc2_
tokens))

# Print the Jaccard Simialrity score
print(jaccard_similarity)

# I like dogs hate
# s1: I like dogs.->1 1 1 0 (one hot encoding)
# s2: I hate dogs->1 0 1 1

from scipy.spatial import distance
distance.dice([1,1,1,0],[1,0,1,1])

```

#Cosine Similarity using Spacy

```

import en_core_web_sm

nlp = en_core_web_sm.load()

## try medium or large spacy english models

doc1 = nlp("I like apples.")
doc2 = nlp("I like oranges.")

# cosine similarity
doc1.similarity(doc2)

doc1 = nlp("I like dogs.")

```

```
doc2 = nlp("I hate dogs.")
```

```
# cosine similarity  
doc1.similarity(doc2)
```

#Cosine Similarity using Scipy

```
from scipy import spatial  
  
# Document Vectorization  
doc1, doc2 = nlp('I like apples.').vector, nlp('I like  
oranges.').vector  
print(doc1)  
print(type(doc1))  
print(doc1.shape)  
  
print(doc2)  
  
# Cosine Similarity  
result = 1 - spatial.distance.cosine(doc1, doc2)  
  
print(result)
```

#Let's create a search engine using Text Similarity measures

```
import en_core_web_sm  
from numpy import dot  
from numpy.linalg import norm  
import numpy as np  
import pandas as pd  
  
nlp = en_core_web_sm.load()  
  
# Prepare dataset  
doc_list=['I love this sandwich.',  
          'this is an amazing place!',  
          'I feel very good about these beers.',  
          'this is my best work.',  
          'what an awesome view',  
          'I do not like this restaurant',  
          'I am tired of this stuff.',  
          "I can't deal with this",  
          'he is my sworn enemy!',  
          'my boss is horrible.',  
          'I hate this sandwich.']  
  
# user input  
query=input()
```

```

# similarity score
sim_scores=[]

# Vectorize input query
q_vector=nlp(query).vector

for doc in doc_list:
    # Vectorize document
    doc_vector=nlp(doc).vector
    # Cosine Similarity
    cos_sim = dot(q_vector,
doc_vector)/(norm(q_vector)*norm(doc_vector))
    # append the score
    sim_scores.append(cos_sim)

# most similar
most_similar=doc_list[sim_scores.index(max(sim_scores))]
print("\nMost Similar:\n",most_similar)

# sorting most similar sentences
top_index=list(np.argsort(sim_scores)[-5:]) #argsort() by default
sorts in ascending order , therefore this part selects the last 5
indices from the sorted list. Since the indices are sorted in
ascending order of similarity scores, selecting the last 5 indices
effectively gives us the indices corresponding to the top 5 most
similar sentences.

# [0.8, 0.4, 0.9, 0.45]
# argsort: [0.4,0.45, 0.8, 0.9]-> 1,3,0,2

print(top_index)
top_index.reverse()
print(top_index)

print("\nMost Similar Documents:\n")
for i in top_index:
    print(doc_list[i],sim_scores[i])

print(sim_scores)

```

#Training Word2vec model

```
import re # For preprocessing
import pandas as pd # For data handling (alias)
from time import time # To time our operations
from collections import defaultdict # For word frequency

import spacy # For preprocessing

import logging # Setting up the loggings to monitor gensim
logging.basicConfig(format="%(levelname)s - %(asctime)s: %(message)s",
                    datefmt= '%H:%M:%S', level=logging.INFO)

df=pd.read_csv('/content/simpsons_script_lines.csv')
df.head()

<ipython-input-2-ee713851d592>:1: DtypeWarning: Columns (4,5,6) have
mixed types. Specify dtype option on import or set low_memory=False.
df=pd.read_csv('/content/simpsons_script_lines.csv')

{"type": "dataframe", "variable_name": "df"}

df.columns

Index(['id', 'episode_id', 'number', 'raw_text', 'timestamp_in_ms',
      'speaking_line', 'character_id', 'location_id',
      'raw_character_text',
      'raw_location_text', 'spoken_words', 'normalized_text',
      'word_count'],
      dtype='object')

df=df[['raw_character_text', 'spoken_words']]

df

{"type": "dataframe", "variable_name": "df"}

df.isnull().sum()

raw_character_text    17522
spoken_words          26159
dtype: int64

# nlp = spacy.load('en', disable=['ner', 'parser'])
nlp = spacy.load("en_core_web_sm")

def cleaning(doc):
    # Lemmatizes and removes stopwords
    # doc needs to be a spacy Doc object
    txt = [token.lemma_ for token in doc if not token.is_stop]
```

```

        if len(txt) > 2:
            return ' '.join(txt)

brief_cleaning = (re.sub("[^A-Za-z']+", ' ', str(row)).lower() for row
in df['spoken_words'])

# t = time()
# txt = [cleaning(doc) for doc in nlp.pipe(brief_cleaning,
batch_size=5000,
#                                     n_threads=-1)]
# print('Time to clean up everything: {} mins'.format(round((time() -
t) / 60, 2)))
t = time()
txt = [cleaning(doc) for doc in nlp.pipe(brief_cleaning,
batch_size=5000)]
print('Time to clean up everything: {} mins'.format(round((time() - t)
/ 60, 2)))

Time to clean up everything: 4.09 mins

df_clean = pd.DataFrame({'clean': txt})
df_clean = df_clean.dropna().drop_duplicates()
df_clean.shape

(86211, 1)

from gensim.models.phrases import Phrases, Phraser
sent = [row.split() for row in df_clean['clean']]

phrases = Phrases(sent, min_count=30, progress_per=10000)

phrases

<gensim.models.phrases.Phrases at 0x79b0b11c0430>

bigram = Phraser(phrases)

# I love ice cream

sentences = bigram[sent]

word_freq = defaultdict(int)
for sent in sentences:
    for i in sent:
        word_freq[i] += 1
len(word_freq)

29791

sorted(word_freq, key=word_freq.get, reverse=True)[:10]

```

```
['oh', 'like', 'know', 'get', 'hey', 'think', 'come', 'right', 'look', 'want']
```

```
import multiprocessing
```

```
from gensim.models import Word2Vec
```

```
cores = multiprocessing.cpu_count() # Count the number of cores in a computer
```

```
w2v_model = Word2Vec(min_count=20,  
                     window=2,  
                     vector_size=300,  
                     sample=6e-5,  
                     alpha=0.03,  
                     min_alpha=0.0007,  
                     negative=20,  
                     workers=cores-1)
```

```
t = time()
```

```
w2v_model.
```

```
print('Time to build vocab: {} mins'.format(round((time() - t) / 60,  
2)))
```

```
Time to build vocab: 0.01 mins
```

```
I like dogs and cats and rabbits  
1 2      3 4 5      4 6
```

```
{I: 1, like: 2..}
```

```
t = time()
```

```
# print(total_examples)
```

```
w2v_model.train(sentences, total_examples=w2v_model.corpus_count,  
epochs=30, report_delay=1)
```

```
print('Time to train the model: {} mins'.format(round((time() - t) /  
60, 2)))
```

```
Time to train the model: 1.39 mins
```

```
w2v_model.corpus_count
```

```
86211
```

```
sentences
```

```
<gensim.interfaces.TransformedCorpus at 0x79b0a94e8130>
```

```
w2v_model.init_sims(replace=True)
```

```
<ipython-input-47-c7757d71a30b>:1: DeprecationWarning: Call to deprecated `init_sims` (Gensim 4.0.0 implemented internal optimizations that make calls to init_sims() unnecessary. init_sims() is now obsolete and will be completely removed in future versions. See https://github.com/RaRe-Technologies/gensim/wiki/Migrating-from-Gensim-3.x-to-4).
```

```
w2v_model.init_sims(replace=True)
```

```
WARNING:gensim.models.keyedvectors:destructive init_sims(replace=True) deprecated & no longer required for space-efficiency
```

```
similar_words = w2v_model.wv.most_similar(positive=["homer"])
for word, similarity in similar_words:
    print(f"{word}: {similarity}")
```

```
gee: 0.8816429376602173
quick: 0.8771034479141235
sweetheart: 0.867999792098999
marge: 0.8667840957641602
hopeless: 0.862589418888092
glad: 0.8622774481773376
bartender: 0.8606762886047363
right: 0.858879029750824
sorry: 0.8574423789978027
help: 0.8572109937667847
```

```
w2v_model.wv.most_similar(positive=["homer_simpson"])
```

```
[('select', 0.9020684361457825),
 ('lady_gentleman', 0.8911157846450806),
 ('trial', 0.887400209903717),
 ('dedicate', 0.8839735388755798),
 ('crew', 0.8823495507240295),
 ('congratulation', 0.8809747695922852),
 ('broadcast', 0.8807786703109741),
 ('convention', 0.879878044128418),
 ('direct', 0.8789938688278198),
 ('host', 0.8774415254592896)]
```

```
w2v_model.wv.most_similar(positive=["marge"])
```

```
[('glad', 0.8914166688919067),
 ('sure', 0.8864542245864868),
 ('pregnant', 0.8832345008850098),
 ('homie', 0.8820032477378845),
 ('talk', 0.8817054033279419),
 ('advice', 0.8815292716026306),
 ('worried', 0.8806873559951782),
 ('someday', 0.8791307806968689),
```

```

('gee', 0.8790706396102905),
('darling', 0.8790403604507446)]

w2v_model.wv.most_similar(positive=["bart"])

[('mom_dad', 0.9176613688468933),
 ('lisa', 0.9168039560317993),
 ('upset', 0.9008905291557312),
 ('worried', 0.9008179903030396),
 ('mother', 0.898597776889801),
 ('concerned', 0.8945304751396179),
 ('admit', 0.8920857906341553),
 ('brother', 0.8900191783905029),
 ('worry', 0.8900096416473389),
 ('humiliate', 0.8885989785194397)]

w2v_model.wv.similarity('maggie', 'baby')
0.7319323

w2v_model.wv.similarity('bart', 'nelson')
0.81326276

w2v_model.wv.doesnt_match(['jimbo', 'milhouse', 'kearney'])
WARNING:gensim.models.keyedvectors:vectors for words {'kearney'} are
not present in the model, ignoring these words

{"type": "string"}

w2v_model.wv.doesnt_match(['homer', 'patty', 'selma'])

{"type": "string"}

w2v_model.wv.most_similar(positive=["woman", "homer"],
negative=["marge"], topn=3)

[('screw', 0.844830334186554),
 ('effort', 0.8303388357162476),
 ('attractive', 0.8162267208099365)]

w2v_model.wv.most_similar(positive=["woman", "bart"],
negative=["man"], topn=3)

[('important', 0.821419358253479),
 ('understand', 0.809717059135437),
 ('know', 0.8067172765731812)]

w2v_model.wv.most_similar(positive=["woman", "queen"],
negative=["man"], topn=3)

```



```
[('spelling', 0.5295670628547668),  
 ('king', 0.5204694271087646),  
 ('pageant', 0.5008288025856018)]
```

```
# w2v_model.wv.most_similar(a=["woman", "queen"], b=["man"], topn=3)
```

```
-----
```

```
-----  
TypeError                                Traceback (most recent call  
last)
```

```
<ipython-input-31-e3aac44ae70c> in <cell line: 1>()
```

```
----> 1 w2v_model.wv.most_similar(a=["woman", "queen"], b=["man"],  
topn=3)
```

```
TypeError: KeyedVectors.most_similar() got an unexpected keyword  
argument 'a'
```